

Proposed Architecture for the K-Square Programme Onboarding Agent

Alberto Espinosa
KSquare Group

1 August 2025

Abstract

The K-Square Programme Onboarding Agent streamlines the onboarding of programme execution teams by automating data aggregation and insight generation. This document proposes a robust architecture to support this goal, detailing agentic separation, system layers, and recommended technologies. The architecture leverages LangGraph for agent orchestration, Ollama for local language model inference, and a conversational AI interface with a visual dashboard for user interaction. Four diagrams illustrate the system's design, providing clarity on agent interactions, data flow, architecture, and high-level workflow.

1 Introduction

The K-Square Programme Onboarding Agent addresses the challenge of manual onboarding by providing execution teams with rapid access to client profiles, domain knowledge, and actionable insights within 2–3 days. A well-designed architecture is essential to ensure seamless data processing, agent coordination, and user interaction via a conversational interface and visual dashboard. This proposal outlines a layered architecture with six specialised components, focusing on technical feasibility and alignment with enterprise workflows.

2 Architecture Overview

The architecture is structured into five layers: Data, Agent, Processing, UI, and Orchestration. Six components operate within the system: Programme Setup Agent, Domain Knowledge Agent, Client Profile Agent, Actionable Insights Agent, Meetings Agent, and Knowledge Base. The system integrates internal data from repositories (e.g., SharePoint) and external data from public web sources, delivering outputs through a conversational AI and a visual dashboard. LangGraph orchestrates agent workflows, while Ollama powers local inference for privacy-sensitive tasks.

3 Agentic Separation

The system employs five distinct agents and a central Knowledge Base, each with specific roles, inputs, and outputs:

3.1 Programme Setup Agent

The Programme Setup Agent guides users through a conversational interface to input project and client details, such as client name, industry, problem statement, and SharePoint links. It queries internal repositories (e.g., SharePoint) and external web sources, validates data sufficiency with users, and routes validated data to the Domain Knowledge and Client Profile Agents. Outputs include structured project profiles stored in the Knowledge Base.

3.2 Domain Knowledge Agent

The Domain Knowledge Agent aggregates internal and external data to provide domain-specific insights, such as industry best practices and reusable templates. Inputs include validated data from the Programme Setup Agent; outputs include tagged summaries (e.g., definitions, lessons learned) stored in the Knowledge Base, feeding into the Actionable Insights Agent.

3.3 Client Profile Agent

The Client Profile Agent builds detailed client profiles using internal documents (e.g., statements of work) and public data (e.g., company websites, LinkedIn profiles). It incorporates user validation to ensure relevance, producing tagged profile cards (e.g., stakeholder details, market position) for the Knowledge Base.

3.4 Actionable Insights Agent

The Actionable Insights Agent synthesises outputs from the Domain Knowledge, Client Profile, and Meetings Agents to generate recommendations, summaries, and tasks. Outputs include checklists, one-page summaries, and alerts for pending actions, stored in the Knowledge Base and displayed via the visual dashboard.

3.5 Meetings Agent

The Meetings Agent analyses Microsoft Teams recordings and transcripts to extract action items, scope insights, and engagement metrics. It employs speech-to-text and natural language processing to provide insights that enhance client interactions, storing outputs in the Knowledge Base for the Actionable Insights Agent.

3.6 Knowledge Base

The Knowledge Base centralises tagged data from all agents, including client profiles, domain knowledge, and recommendations. It supports a tagging system for organisation (e.g., Client, Problem Statement, Best Practices) and enables search functionality for internal and external data, feeding the visual dashboard.

4 Layer Descriptions

The architecture comprises five layers, each addressing a critical aspect of the system:

4.1 Data Layer

The Data Layer manages ingestion, storage, and preprocessing of internal and external data. Internal data (e.g., statements of work, meeting recordings) are retrieved from SharePoint via the Microsoft Graph API. External data (e.g., industry reports, public records) are collected through web crawling. A vector database (e.g., Pinecone) stores document embeddings for semantic search. Preprocessing involves text extraction using optical character recognition (e.g., Tesseract) and chunking with natural language processing libraries (e.g., SpaCy).

4.2 Agent Layer

The Agent Layer hosts the Programme Setup, Domain Knowledge, Client Profile, Actionable Insights, and Meetings Agents, coordinated by LangGraph. The Programme Setup Agent initiates data collection via conversational inputs, routing data to other agents. The Knowledge Base stores tagged outputs. Ollama runs local large language models (e.g., LLaMA 3) for internal data processing, with cloud-based models (e.g., OpenAI) for external data.

4.3 Processing Layer

The Processing Layer handles language model inference, vector search, and feedback integration. Large language models are fine-tuned using techniques like LoRA for domain-specific accuracy. Vector search, powered by Pinecone or Weaviate, enables rapid document retrieval. A feedback mechanism collects user validations (e.g., relevance scores) to retrain models, targeting 60

4.4 UI Layer

The UI Layer provides a React-based dashboard with tabs for client profiles, domain knowledge, meeting insights, and recommendations. It features a conversational interface for data input, visual cards for tagged data, and action buttons (e.g., "Generate Checklist"). Chart.js powers visualisations (e.g., engagement metrics), and a REST API (e.g., FastAPI) connects the dashboard to the backend.

4.5 Orchestration Layer

The Orchestration Layer, powered by LangGraph, manages agent workflows and system monitoring. Workflows are defined as graphs, with nodes representing agent tasks and edges denoting data flow to the Knowledge Base. Monitoring tools (e.g., Prometheus, Grafana) track agent performance and language model accuracy, ensuring reliability and scalability.

5 Proposed Technologies

The following technologies are recommended, with justifications and alternatives:

5.1 LangGraph

LangGraph orchestrates agent workflows, enabling dynamic task routing and context retention. It supports complex interactions, such as routing Programme Setup Agent outputs to Domain Knowledge and Client Profile Agents. Its complexity requires experienced developers, but CrewAI is a simpler alternative for prototyping, though less scalable.

5.2 Ollama

Ollama enables local language model inference (e.g., LLaMA 3) for processing internal SharePoint documents, ensuring data privacy. Its limitations in large-scale inference necessitate cloud-based models (e.g., OpenAI, Google Gemini) for external data. A hybrid approach balances privacy and performance.

5.3 Supporting Technologies

- **Pinecone:** A vector database for semantic search, offering scalability. Weaviate is an open-source alternative.
- **Whisper:** OpenAI's speech-to-text model for transcribing meeting recordings, with high accuracy across languages.
- **FastAPI:** A lightweight framework for building REST APIs to connect the UI and backend.
- **React:** A robust library for an interactive dashboard, styled with Tailwind CSS for responsiveness.
- **Prometheus and Grafana:** Monitoring tools to ensure system reliability and performance.

6 System Diagrams

The system's design is illustrated through four diagrams, each presented on a dedicated landscape page for clarity. These diagrams clarify agent interactions, data flow, architectural layers, and the high-level workflow.

6.1 Agents Interaction Diagram

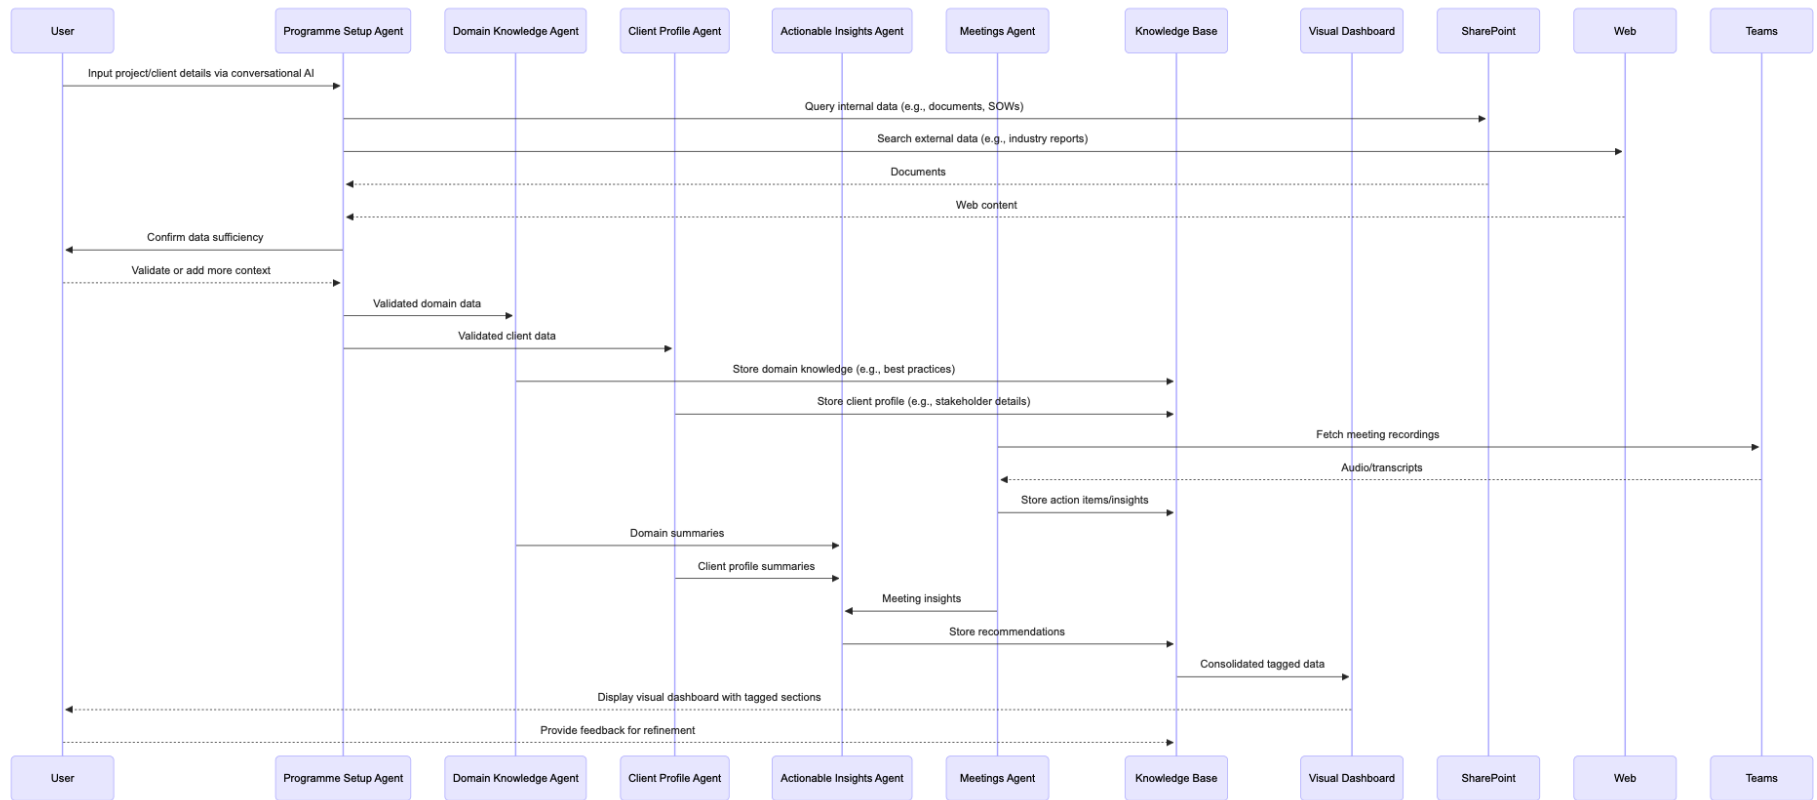


Figure 1: Agents Interaction Diagram, showing conversational AI and visual dashboard interactions.

Figure 1 is a sequence diagram illustrating interactions between the User, Programme Setup Agent, Domain Knowledge Agent, Client Profile Agent, Actionable Insights Agent, Meetings Agent, Knowledge Base, and Visual Dashboard. The process begins with users providing project details via a conversational AI, followed by data queries, user validation, and storage in the Knowledge Base. Outputs are displayed in a visual dashboard, with feedback refining agent performance.

6.2 Data Flow Diagram

Figure 2 depicts data flow from SharePoint and web sources to the vector database, Programme Setup Agent, and other agents. Validated outputs are stored in the Knowledge Base, feeding the Actionable Insights Agent and Visual Dashboard via a REST API. Feedback loops refine agent outputs, ensuring a closed-loop system.

6.3 Architecture Diagram

Figure 3 outlines the five layers: Data, Agent, Processing, UI, and Orchestration. The Agent Layer includes the Programme Setup, Domain Knowledge, Client Profile, Actionable Insights, and Meetings Agents, with the Knowledge Base centralising data. LangGraph orchestrates workflows, and the React dashboard delivers outputs via a REST API.

6.4 High-Level Overview Diagram

Figure 4 presents a simplified view of the system’s workflow, from client/project inputs through the Data, Agent, Processing, and Orchestration Layers to onboarding outputs in the UI Layer. Feedback loops ensure continuous improvement.

7 Suggested Additional Features

The Meetings Agent utilises speech-to-text and natural language processing to extract insights from meeting data. A novel approach proposed by Espinosa [1] introduces a dynamical measurement of opinion change based on semantics and complexity, detailed at arXiv:2505.02581. This method uses perturbation and intervention analysis to track conversational dynamics, enabling the Meetings Agent to identify nuanced client perspective shifts. Integrating this feature could enhance insight precision, subject to evaluation for feasibility.

8 Conclusion

The proposed architecture for the K-Square Programme Onboarding Agent provides a scalable, privacy-conscious solution for automating team onboarding. By leveraging LangGraph, Ollama, and supporting technologies like Pinecone and Whisper, the system ensures robust performance. The conversational AI, visual dashboard, and Knowledge Base enhance usability, while the layered design supports future enhancements, such as advanced conversational analysis.

References

- [1] A. Espinosa, “Embracing Inevitable AI Misalignment as a Strategy to Steer Competitive Agents Towards Human Alignment: Change-of-Opinion Attacks and Interventions,” arXiv preprint arXiv:2505.02581, 2025.

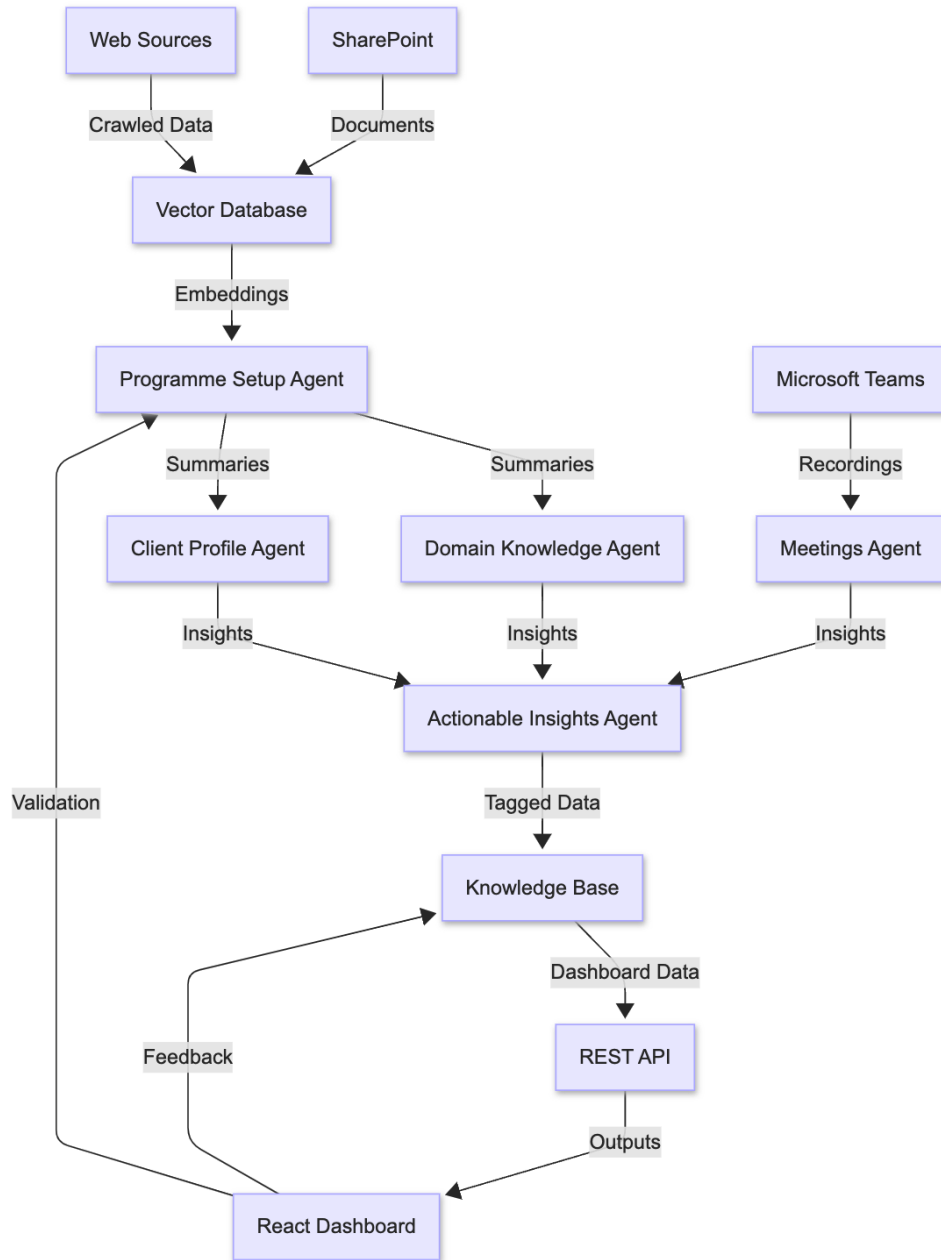


Figure 2: Data Flow Diagram, illustrating data movement through agents and the Knowledge Base.

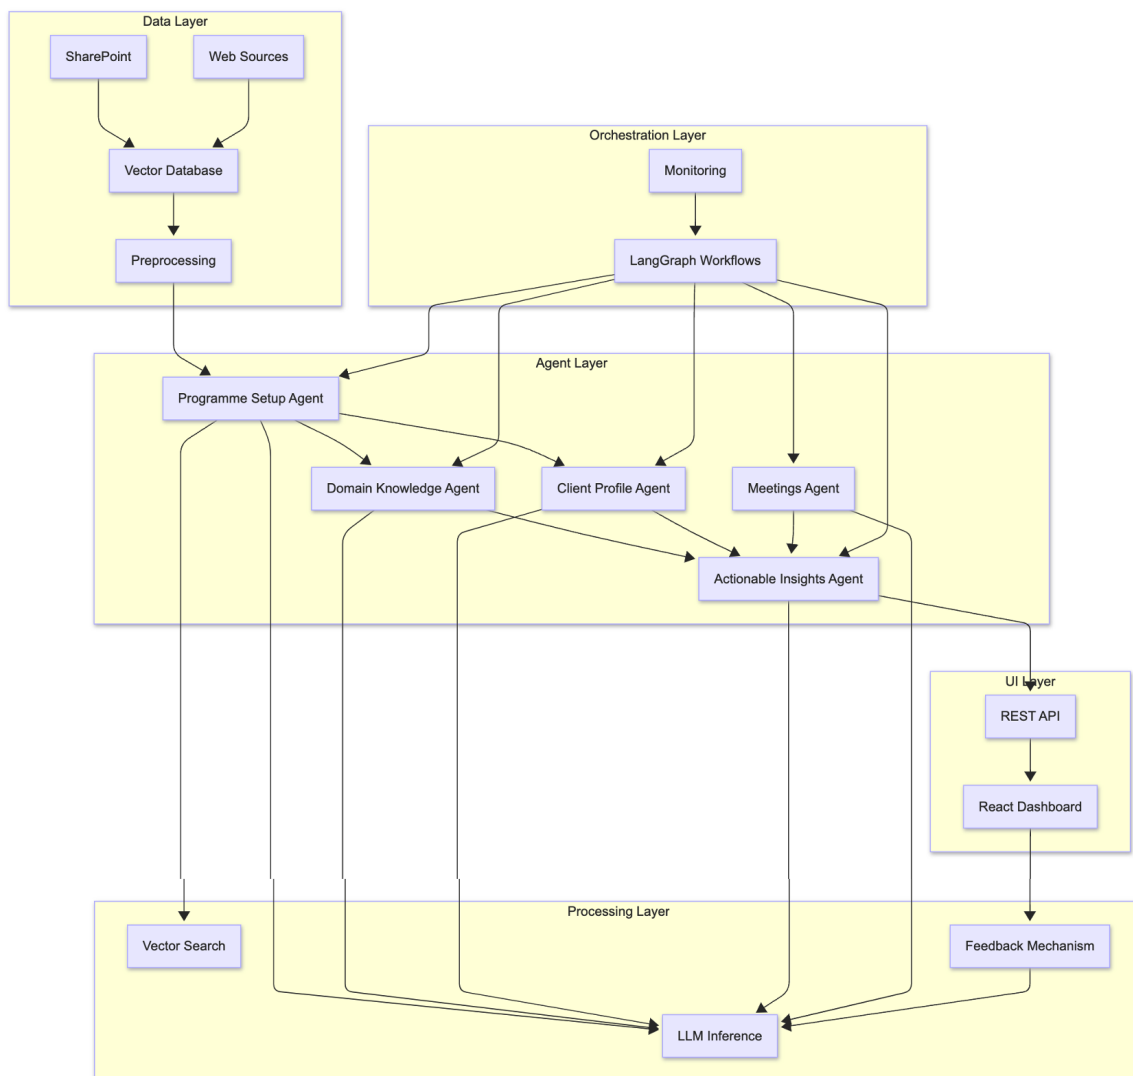


Figure 3: Architecture Diagram, showing the five-layer structure and agent interactions.




Figure 4: High-Level Overview Diagram, showing the end-to-end workflow.